

ezr-w2.lua

ezr, week 2: tables, distance, gap to heaven. One weekly-sized chunk of [ezr.lua](#), verbatim, with glossary notes folded in, and this week’s exercises at the bottom.

This week’s story

Rows gather into tables; row 1 alone decides every column’s kind and role. Then two numbers see the whole table: distx, the gap between two rows over the x columns; and disty, the gap from a row’s goals to heaven. Sorting by disty is optimization with no model, no weights, no training.

▼ Tbl

Rows, plus the column summaries those rows built. Row 1 of any source is the header, and the header alone decides each column’s kind ([noir](#)) and role: a trailing [!](#) is the class, [+](#) or [-](#) a goal (a y column), [X](#) is ignored, the rest are the x (independent) columns.

```
function Tbl(src)
  src = iter(src)
  return adds(src, new(TBL, {rows={}, mid=nil,
                        cols=Cols(src())})) end

function Cols(names, all,x,y,klass)
  all, x, y = {}, {}, {}
  for at, s in ipairs(names) do
    all[at] = Col(s, at)
    if s:find"!$" then klass = all[at]
    elseif s:find"[+-$]" then y[#y+1] = all[at]
    elseif s:sub(-1) ~= "X" then x[#x+1] = all[at] end end
  return new(COLS, {names=names,all=all,x=x,y=y,klass=class}) end
```

▼ minkowski

Minkowski’s p-norm, folding many per-column gaps g_c (each 0..1) into one 0..1 number:

$$d = ((1)/(n)\sum g_c^{p})^{1/p}$$

$p=1$ is the Manhattan distance (all gaps count equally), $p=2$ Euclidean; as p grows, the largest single gap dominates. [the.p](#) defaults to 2.

```
function minkowski(cols,f, d,n)
  d, n = 0, TINY
  for _, c in ipairs(cols) do n, d = n+1, d + f(c) ^ the.p end
  return (d / n) ^ (1 / the.p) end
```

Note the design: minkowski never touches the data. It takes a function [f](#) and calls [f\(C\)](#) per column, on demand — a higher-order, lazy style. Each caller passes its own little lambda (distx measures row gaps, disty measures gaps to heaven), and no intermediate list of gaps is ever built. The Python analog is a generator expression, computing each term only as it is summed:

```
d = (sum(f(c)**p for c in cols) / len(cols)) ** (1/p)
```

```
function Tbl(src)
  src = iter(src)
  return adds(src, new(TBL, {rows={}, mid=nil,
                        cols=Cols(src())})) end

function Cols(names, all,x,y,klass)
  all, x, y = {}, {}, {}
  for at, s in ipairs(names) do
    all[at] = Col(s, at)
    if s:find"!$" then klass = all[at]
    elseif s:find"[+-$]" then y[#y+1] = all[at]
    elseif s:sub(-1) ~= "X" then x[#x+1] = all[at] end end
  return new(COLS, {names=names,all=all,x=x,y=y,klass=class}) end
```

```
function minkowski(cols,f, d,n)
  d, n = 0, TINY
  for _, c in ipairs(cols) do n, d = n+1, d + f(c) ^ the.p end
  return (d / n) ^ (1 / the.p) end
```

```
function SYM.dist(i,a,b)
  if a == "?" and b == "?" then return 1 end
  return a ~= b and 1 or 0 end

function NUM.dist(i,a,b)
  if a == "?" and b == "?" then return 1 end
  a, b = i:norm(a), i:norm(b)
  if a == "?" then a = b > 0.5 and 0 or 1 end
  if b == "?" then b = a > 0.5 and 0 or 1 end
  return abs(a - b) end

function TBL.distx(i,row1,row2)
  return minkowski(i.cols.x, function(c)
    return c:dist(row1[c.at], row2[c.at]) end) end
```

▼ distx

The gap between two rows, over the x columns only: each column measures its own 0..1 gap (`dist` in the `columnProtocol`), and `minkowski` folds them. An unknown `"?"` assumes the worst: symbols that might differ, do; a missing number is pushed to whichever end is further away.

```
function SYM.dist(i,a,b)
  if a == "?" and b == "?" then return 1 end
  return a ~= b and 1 or 0 end

function NUM.dist(i,a,b)
  if a == "?" and b == "?" then return 1 end
  a, b = i:norm(a), i:norm(b)
  if a == "?" then a = b > 0.5 and 0 or 1 end
  if b == "?" then b = a > 0.5 and 0 or 1 end
  return abs(a - b) end

function TBL.distx(i,row1,row2)
  return minkowski(i.cols.x, function(c)
    return c:dist(row1[c.at], row2[c.at]) end) end
```

▼ heaven

The best value a goal column can hope for, in normalized 0..1 space: 0 for a minimize goal (trailing `-`), 1 for a maximize goal (trailing `+`). Decided in one line, at column birth:

```
heaven = name:find"-$" and 0 or 1
```

▼ disty


The gap from a row's goals to `heaven`: each y column measures $|norm(v) - heaven|$, and `minkowski` folds them. 0 = best possible row, 1 = worst. No model, no weights, no training — sort rows by `disty` and the best float to the top (`--disty` in `ezr-eg2`). `disty` is an `*aggregation*`, the zooming rival to frontier-chasing (see `Pareto zoom effect`). Later weeks make `disty` the thing that costs money: it reads the goal columns, and goals are labels.

```
function TBL.disty(i,row)
  if i.model and row[i.cols.y[1].at] == "?" then
    i:label(row) end
  return minkowski(i.cols.y, function(y)
    return abs(y:norm(row[y.at]) - y.heaven) end) end
```

▼ Pareto frontier

"Give me the fruitful error any time, full of seeds, bursting with its own corrections. You can keep your sterile truth for yourself." — Vilfredo Pareto

With many goals there is rarely one best row — lighter cars brake worse. Row `a` `*dominates*` `b` if `a` is at least as good on every goal and better on one. The frontier (`O`) is whatever nothing dominates:



Report the frontier and let the customer pick their trade-off.

```
function TBL.disty(i,row)
  if i.model and row[i.cols.y[1].at] == "?" then
    i:label(row) end
  return minkowski(i.cols.y, function(y)
    return abs(y:norm(row[y.at]) - y.heaven) end) end
```

▼ Pareto zoom effect

Ganguly & Menzies, "[Zoom, Don't Wander](#)" (2026): across 100+ SE optimization tasks, Pareto-optimal solutions are RARE (about 0.6% of configurations) and CLUMPED — a tiny island, tight in decision space (85% of datasets) and huddled near the ideal corner of objective space (88%). Real example, the Redis configuration landscape (from [PromiseTune](#), Chen & Chen, ICSE 2026): the dark-red good region is a small fraction of a rugged space, and tuners that wander it (SMAC, random search) plateau far below the optimum:



PromiseTune Fig 1: Redis configuration landscape and tuning trajectories

So frontier-chasing evolvers and global Bayesian methods spend most of a small labelling budget wandering the huge ungood region; a greedy regional search that zooms toward the island wins or ties in 84-89% of cases, running 2-3 orders of magnitude faster. The opposite of frontier reasoning is an *aggregation function* — collapse all goals to one number and chase that. [disty](#) is this course's aggregation function: zoom, don't wander.

exercises

****Exercises, week 2**** 1. Port this page's code to Python, wiring each demo of `ezr-eg2 (~distx --disty -laws~)` to a `test_` function. 2. On `auto93`, sort all rows by `disty` and print the top and bottom three. Do the best rows *look* best? (If not, see exercise 5 of week 1.) 3. `--laws~` probes distance laws at random. State the four laws. Which one fails for the "?"-handling here, and why is that a price worth paying? 4. Set `the.p=1`, then 4, then 8. How does the `disty` ranking change? What is `minkowski` converging to as `p` grows? 5. `distx` never reads a goal; `disty` reads nothing else. Why does that split matter when labels cost money?

local _ = "week 2 ends here"